

```
• using Statistics
```

```
• using LinearAlgebra
```

```
• using Plots
```

```
• using PlutoUI
```

```
PlotlyBackend()
```

```
• plotly()
```

For saving to png with the Plotly backend PlotlyBase has to be installed.

Decision Trees!

Decision trees in Julia.

Tree Structure

Defining the base tree structure for the decision tree.

```
• mutable struct DecisionNode
•     threshold::Union{Number, Missing}
•     attribute::Union{Integer, Missing}
•     parent::Union{DecisionNode, Missing}
•     leftChild::Union{DecisionNode, LeafNode, Missing}
•     rightChild::Union{DecisionNode, LeafNode, Missing}
•
•     DecisionNode() = new(missing, missing, missing, missing, missing)
•     DecisionNode(threshold::Number, attribute::Integer) = new(threshold,
attribute, missing, missing, missing)
•     DecisionNode(threshold::Number, attribute::Integer, parent::DecisionNode) =
new(threshold, attribute, parent, missing, missing)
• end
```

```
• mutable struct LeafNode
•     class::Bool
•
•     LeafNode(class::Bool) = new(class)
• end
```

setleftchild!

setleftchild!(node::DecisionNode, child::DecisionNode)

Sets the leftChild of the node to child, if it has not been already set.

```
• """setleftchild!(node::DecisionNode, child::DecisionNode)
•
• Sets the `leftChild` of the `node` to `child`, if it has not been already set.
• """
• function setleftchild!(node::DecisionNode, child::DecisionNode)
•     if !isequal(node.leftChild, missing)
•         error("Left child has already been set!")
•     else
•         node.leftChild = child
•     end
• end
```

setrightchild!

setrightchild!(node::DecisionNode, child::DecisionNode)

Sets the rightChild of the node to child, if it has not been already set.

```
• """setrightchild!(node::DecisionNode, child::DecisionNode)
•
• Sets the `rightChild` of the `node` to `child`, if it has not been already set.
• """
• function setrightchild!(node::DecisionNode, child::DecisionNode)
•     if !isequal(node.rightChild, missing)
•         error("Right child has already been set!")
•     else
•         node.rightChild = child
•     end
• end
```

hasleftchild

hasleftchild(node::DecisionNode)

Checks whether the node has a left child.

```
• """hasleftchild(node::DecisionNode)
•
• Checks whether the node has a left child.
• """
• function hasleftchild(node::DecisionNode)
•     return !isequal(node.leftChild, missing)
• end
```

hasrightchild

hasrightchild(node::DecisionNode)

Checks whether the node has a left child.

```
• """hasrightchild(node::DecisionNode)
•
• Checks whether the node has a left child.
• """
• function hasrightchild(node::DecisionNode)
•     return !isequal(node.rightChild, missing)
• end
```

inorder_rec

inorder_rec(node::DecisionNode)

Recursive function printing the threshold values within the tree while traversing it in order.

```

• """inorder_rec(node::DecisionNode)
•
• Recursive function printing the threshold values within the tree while
• traversing it in order.
• """
• function inorder_rec(node::DecisionNode)
•     if hasleftchild(node)
•         inorder_rec(node.leftChild)
•     end
•     print(node.threshold, " ")
•     if hasrightchild(node)
•         inorder_rec(node.rightChild)
•     end
• end

```

inorder

inorder(node::DecisionNode)

Function printing the threshold values within the tree while traversing it in order.

```

• """inorder(node::DecisionNode)
•
• Function printing the threshold values within the tree while traversing it in
• order.
• """
• function inorder(node::DecisionNode)
•     with_terminal() do
•         inorder_rec(node)
•     end
• end

```

Entropy and Information Gain

entropy

entropy(Y::AbstractArray{Bool})

Computes the Shannon's entropy of a given boolean array.

```

• """entropy(Y::AbstractArray{Bool})
•
• Computes the Shannon's entropy of a given boolean array.
• """
• function entropy(Y::AbstractArray{Bool})
•     positiveProb = sum(Y) / length(Y)
•     negativeProb = 1 - positiveProb
•
•     if length(Y) == 0
•         return 0.0
•     elseif positiveProb == 1 || negativeProb == 1
•         return 0.0
•     end
•
•     return -(positiveProb * log(positiveProb) + negativeProb *
• log(negativeProb))
• end

```

informationgain

informationgain(Y::AbstractArray{Bool}, Y1::AbstractArray{Bool}, Y2::AbstractArray{Bool})

Computes the information gain obtained by dividing the boolean array Y into the two boolean arrays Y1 and Y2.

```

• """informationgain(Y::AbstractArray{Bool}, Y1::AbstractArray{Bool},
• Y2::AbstractArray{Bool})
•
• Computes the information gain obtained by dividing the boolean array `Y` into
• the two boolean arrays `Y1` and `Y2`.
• """
• function informationgain(Y::AbstractArray{Bool}, Y1::AbstractArray{Bool},
• Y2::AbstractArray{Bool})
•     return entropy(Y) - entropy(Y2) - entropy(Y1)
• end

```

Dataset generation

Generating a simple random dataset.

```
md"""### Dataset generation
Generating a simple random dataset.
"""
```

rownorm

`rownorm(X::AbstractMatrix)`

Computes the norms of vectors forming the rows of the matrix `X`.

```
"""rownorm(X::AbstractMatrix)

Computes the norms of vectors forming the rows of the matrix `X`.
"""
function rownorm(X::AbstractMatrix)
    rowCount = size(X, 1)
    return [ norm(X[i, :]) for i in 1:rowCount ]
end
```

generatedataset

generatedataset(examplesCount::Integer)

Generates a simple dataset.

```

• """generatedataset(examplesCount::Integer)
•
• Generates a simple dataset.
• """
• function generatedataset(examplesCount::Integer)
•     positiveCount = div(examplesCount, 2)
•     negativeCount = examplesCount - positiveCount
•
•     positiveExamples = randn(positiveCount, 2)
•     negativeExamples = randn(negativeCount, 2) .* 4
•
•     minNegativeNorm = maximum(rownorm(positiveExamples)) * 1.2
•
•     negativeExamples = negativeExamples[filter(i -> norm(negativeExamples[i,
• :]) > minNegativeNorm, 1:negativeCount), :]
•
•     while size(negativeExamples, 1) < negativeCount
•
•         newNegativeExamples = randn(negativeCount, 2) .* 4
•
•         newNegativeExamples = newNegativeExamples[filter(i ->
norm(newNegativeExamples[i, :]) > minNegativeNorm, 1:negativeCount), :]
•
•         negativeExamples = vcat(negativeExamples, newNegativeExamples)
•     end
•
•     negativeExamples = negativeExamples[1:negativeCount, :]
•
•     X = vcat(positiveExamples, negativeExamples)
•     Y = [ trues(positiveCount); falses(negativeCount) ]
•
•     return X, Y
• end

```

- `X, Y = generatedataset(1000)`

```
plot_examples(X::AbstractMatrix, Y::BitVector)
```

A scatter plot showing two classes, y_1 (blue dots) and y_2 (red dots), in a 2D space. The x-axis ranges from -15 to 15, and the y-axis ranges from -10 to 15. The blue dots are clustered in the center, around (0, 0). The red dots are more widely distributed, forming a ring-like shape around the blue cluster, with some outliers further away. The plot illustrates a non-linearly separable dataset where the two classes are not easily distinguished by a straight line.

- `plot_examples(X, Y)`

Building the Trees

```

• mutable struct DecisionTree
•     root::DecisionNode
•
•     function DecisionTree(root::DecisionNode)
•         new(root)
•     end
• end

```

splitonvalue

splitonvalue(X::AbstractVector, Y::BitVector, value::Number)

Splits the vector `Y` into two parts, `first` and `second`, based on the provided `value` and the vector `X`.

```

• """splitonvalue(X::AbstractVector, Y::BitVector, value::Number)
•
• Splits the vector `Y` into two parts, `first` and `second`, based on the
• provided `value` and the vector `X`.
• """
• function splitonvalue(X::AbstractVector, Y::BitVector, value::Number)
•     first = [ Y[i] for i in eachindex(X) if X[i] <= value ]
•     second = [ Y[i] for i in eachindex(X) if X[i] > value ]
•     return first, second
• end

```

splitonvalue

`splitonvalue(X::AbstractVector, Y::BitVector, value::Number)`

Splits the vector `Y` into two parts, `first` and `second`, based on the provided `value` and the vector `X`.

`splitonvalue(X::AbstractMatrix, Y::BitVector, attrib::Integer, value::Number)`

Splits the dataset `(X, Y)` into two parts, `first` and `second`, based on the provided `value` and `attrib` indicating the column of the matrix `X`.

```

• """splitonvalue(X::AbstractMatrix, Y::BitVector, attrib::Integer, value::Number)
•
• Splits the dataset `(X, Y)` into two parts, `first` and `second`, based on the
• provided `value` and `attrib` indicating the column of the matrix `X`.
• """
• function splitonvalue(X::AbstractMatrix, Y::BitVector, attrib::Integer,
• value::Number)
•     firstX = X[filter(i -> X[i, attrib] <= value, 1:size(X, 1)), :]
•     secondX = X[filter(i -> X[i, attrib] > value, 1:size(X, 1)), :]
•     firstY = Y[filter(i -> X[i, attrib] <= value, 1:size(X, 1))]
•     secondY = Y[filter(i -> X[i, attrib] > value, 1:size(X, 1))]
•     return firstX, secondX, firstY, secondY
• end

```

findbestsplitvalue

findbestsplitvalue(X::AbstractVector, Y::BitVector)

Finds the split on value that yields the largest information gain.

```
• """findbestsplitvalue(X::AbstractVector, Y::BitVector)
•
• Finds the split on value that yields the largest information gain.
• """
• function findbestsplitvalue(X::AbstractVector, Y::BitVector)
•     bestSplitValue = X[1]
•     bestInformationGain = -Inf
•     for value in X
•         first, second = splitonvalue(X, Y, value)
•         informationGain = informationgain(Y, first, second)
•         if informationGain > bestInformationGain
•             bestSplitValue = value
•             bestInformationGain = informationGain
•         end
•     end
•     return bestSplitValue, bestInformationGain
• end
```

findbestsplit

findbestsplit(X::AbstractMatrix, Y::BitVector)

Finds the best attribute and its value on which to split the dataset (X, Y) in order to obtain to largest information gain.

```
• """findbestsplit(X::AbstractMatrix, Y::BitVector)
•
• Finds the best attribute and its value on which to split the dataset `(X, Y)`
• in order to obtain to largest information gain.
• """
• function findbestsplit(X::AbstractMatrix, Y::BitVector)
•     bestSplitAttrib = 0
•     bestSplitValue::Number = 0.0
•     bestInformationGain = -Inf
•     for i in 1:size(X, 2)
•         splitValue, informationGain = findbestsplitvalue(X[:, i], Y)
•         if informationGain > bestInformationGain
•             bestSplitAttrib = i
•             bestSplitValue = splitValue
•             bestInformationGain = informationGain
•         end
•     end
•     return bestSplitAttrib, bestSplitValue
• end
```

splitnode

splitnode(node::DecisionNode, X::AbstractMatrix, Y::BitVector, depth::Integer, maxDepth::Number)

Constructs the decision node `node`, so that it makes such a decision that yields the largest information gain. Recursively constructs its children until no decisions can be made or until the maximal depth of the tree has been reached.

```

"""splitnode(node::DecisionNode, X::AbstractMatrix, Y::BitVector,
depth::Integer, maxDepth::Number)
.
.
Constructs the decision node `node`, so that it makes such a decision that
yields the largest information gain. Recursively constructs its children until
no decisions can be made or until the maximal depth of the tree has been
. reached.
. """
function splitnode(node::DecisionNode, X::AbstractMatrix, Y::BitVector,
. depth::Integer, maxDepth::Number)
.     class = mean(Y) > 0.5
.     if depth >= maxDepth
.         return LeafNode(class)
.     elseif entropy(Y) == 0
.         return LeafNode(class)
.     end
.
.     splitAttrib, splitValue = findbestsplit(X, Y)
.     firstX, secondX, firstY, secondY = splitonvalue(X, Y, splitAttrib,
. splitValue)
.
.     node.attribute = splitAttrib
.     node.threshold = splitValue
.
.     leftChild = DecisionNode()
.     leftChild.parent = node
.     leftChild = splitnode(leftChild, firstX, firstY, depth + 1, maxDepth)
.
.     rightChild = DecisionNode()
.     rightChild.parent = node
.     rightChild = splitnode(rightChild, secondX, secondY, depth + 1, maxDepth)
.
.     node.leftChild = leftChild
.     node.rightChild = rightChild
.
.     return node
.
end

```

decisiontree

decisiontree(X::AbstractMatrix, Y::BitVector, maxDepth::Number = Inf)

Constructs a decision tree based on the given dataset (X, Y).

```

• """decisiontree(X::AbstractMatrix, Y::BitVector, maxDepth::Number = Inf)
•
• Constructs a decision tree based on the given dataset `(X, Y)`.
• """
• function decisiontree(X::AbstractMatrix, Y::BitVector, maxDepth::Number = Inf)
•     root = DecisionNode()
•     splitnode(root, X, Y, 0, maxDepth)
•     return DecisionTree(root)
• end

```

Replacing docs for `Main.workspace#17.decisiontree :: Union{Tuple{AbstractMatrix, BitVector}, Tuple{AbstractMatrix, BitVector, Number}}` in module `Main.workspace#17`

Classification using the Decision Tree

decisiontree_classify

decisiontree_classify(decisionTree::DecisionTree, X::AbstractVector)

Classifies the observation X using the decision tree decisionTree.

```

• """decisiontree_classify(decisionTree::DecisionTree, X::AbstractVector)
•
• Classifies the observation `X` using the decision tree `decisionTree`.
• """
• function decisiontree_classify(decisionTree::DecisionTree, X::AbstractVector)
•     currentNode::Union{DecisionNode, LeafNode} = decisionTree.root
•     while isequal(typeof(currentNode), DecisionNode)
•         currentNode = X[currentNode.attribute] <= currentNode.threshold ?
currentNode.leftChild : currentNode.rightChild
•     end
•
•     return currentNode.class
• end

```

decisiontree_classify

`decisiontree_classify(decisionTree::DecisionTree, X::AbstractVector)`

Classifies the observation `X` using the decision tree `decisionTree`.

`decisiontree_classify(decisionTree::DecisionTree, X::AbstractVector)`

Classifies all of the observations contained in the matrix `X` using the decision tree `decisionTree`.

```

• """decisiontree_classify(decisionTree::DecisionTree, X::AbstractVector)
•
• Classifies all of the observations contained in the matrix `X` using the
• decision tree `decisionTree`.
• """
• function decisiontree_classify(decisionTree::DecisionTree, X::AbstractMatrix)
•     return [ decisiontree_classify(decisionTree, X[i, :]) for i in 1:size(X, 1)
• ]
• end

```

accuracy

`accuracy(expected::AbstractVector{Bool}, predicted::AbstractVector{Bool})`

Computes the classification accuracy given the `expected` and `predicted` vectors.

```

• """accuracy(expected::AbstractVector{Bool}, predicted::AbstractVector{Bool})
•
• Computes the classification accuracy given the `expected` and `predicted`
• vectors.
• """
• function accuracy(expected::AbstractVector{Bool},
• predicted::AbstractVector{Bool})
•     return mean(expected .== predicted)
• end

```

datasetaccuracy

datasetaccuracy(decisionTree::DecisionTree, X::AbstractMatrix, Y::AbstractVector{Bool})

Computes the accuracy of the decision tree `decisionTree` on the dataset `(X, Y)`.

```

• """datasetaccuracy(decisionTree::DecisionTree, X::AbstractMatrix,
  Y::AbstractVector{Bool})
•
• Computes the accuracy of the decision tree `decisionTree` on the dataset `(X,
  Y)`.
• """
• function datasetaccuracy(decisionTree::DecisionTree, X::AbstractMatrix,
  Y::AbstractVector{Bool})
•     return mean(decisiontree_classify(decisionTree, X) .== Y)
• end

```

Check the Accuracy on the Training Dataset

```

basic_tree =
    DecisionTree(DecisionNode(2.638953901238894, 1, missing, DecisionNode(-3.2179346619

```

```

• basic_tree = decisiontree(X, Y, Inf)

```

```

1.0

```

```

• datasetaccuracy(basic_tree, X, Y)

```

Training and Testing datasets

```

• md"""### Training and Testing datasets"

```

```

datasetSize = 1000

```

```

• datasetSize = 1000

```

```

trainDatasetFraction = 0.7

```

```

• trainDatasetFraction = 0.7

```


- datasetX, datasetY = generatedataset(datasetSize)



```
begin
  trainMask = rand(datasetSize) .< trainDatasetFraction
  testMask = map(x -> !x, trainMask)
  trainX = datasetX[trainMask, :]
  trainY = datasetY[trainMask]
  testX = datasetX[testMask, :]
  testY = datasetY[testMask]
end
```

Testing the Tree

tree =

DecisionTree(DecisionNode(2.55241689484124, 1, missing, DecisionNode(-3.15711262695,

- tree = `decisiontree(trainX, trainY, 4)`

1.0

- `datasetaccuracy(tree, trainX, trainY)`

0.9932885906040269

- `datasetaccuracy(tree, testX, testY)`